

# COMP5541 Assignment 1 - Question 2 Results: Convolutional Neural Networks


This document contains the results and analysis from implementing various CNN architectures for classifying images from the CIFAR-10 dataset.

## Part (a): CNN Architecture Comparison

### 1. AlexNet Results

**Architecture:** Simplified AlexNet adapted for CIFAR-10 - 5 convolutional layers + 3 fully connected layers - Dropout regularization in classifier - **Training Setup:** SGD optimizer, lr=0.01, momentum=0.9, Cosine Annealing scheduler


**Performance Summary:** - **Best Test Accuracy:** ~85.7% (achieved around epoch 24) - **Training Characteristics:** - Fast initial convergence - Shows signs of overfitting after epoch 20 - Training accuracy continues improving while test accuracy plateaus



### 2. VGG16 Results

**Architecture:** VGG16 with batch normalization - 13 convolutional layers + 3 fully connected layers - Batch normalization after each convolution - **Training Setup:** Same as AlexNet


**Performance Summary:** - **Best Test Accuracy:** ~90.2% (achieved around epoch 23) - **Training Characteristics:** - Slower initial training due to deeper architecture - Better final performance than AlexNet - More stable training curves with batch normalization



### 3. ResNet18 Results

**Architecture:** ResNet18 with residual connections - 18 layers with skip connections - Batch normalization throughout - **Training Setup:** Same as others

**Performance Summary:** - **Best Test Accuracy:** ~93.2% (achieved around epoch 24) - **Training Characteristics:** - **Best performance** among the three architectures - Most stable training with least overfitting - Residual connections enable effective deep network training -



#### Part (a) Analysis Summary

**Performance Ranking:** ResNet18 > VGG16 > AlexNet

**Key Observations:** 1. **ResNet18** achieved the highest accuracy (93.2%) with the most stable training 2. **VGG16** showed good performance (90.2%) but was computationally heavy 3. **AlexNet** had fastest training but lowest final accuracy (85.7%) with overfitting issues 4. **Residual connections** in ResNet18 proved crucial for training deeper networks effectively 5. **Batch normalization** in VGG16 and ResNet18 helped stabilize training

#### Part (b): Optimizer Comparison on AlexNet

**Training Results with Different Optimizers**

##### 1. SGD with Momentum

- **Learning Rate:** 0.01, Momentum: 0.9
- **Best Test Accuracy:** ~89.7% (epoch 21)

- **Training Pattern:** Steady convergence, best final performance

## 2. RMSProp

- **Learning Rate:** 0.001, Alpha: 0.99
- **Best Test Accuracy:** ~79.4% (epoch 21-23)
- **Training Pattern:** Fast initial convergence, plateaus at lower accuracy




Figure 1: image

•

## 3. Adam

- **Learning Rate:** 0.001, Betas: (0.9, 0.999)
- **Best Test Accuracy:** ~80.4% (epoch 21)
- **Training Pattern:** Very smooth training, stable but moderate performance




Figure 2: image

•

## Part (b) Analysis Summary

**Performance Ranking:** SGD > Adam > RMSProp

**Key Findings:** 1. **SGD with momentum** achieved the best final performance despite slower initial progress 2. **Adaptive optimizers** (Adam, RMSProp) showed faster initial convergence but lower final accuracy 3. **SGD** demonstrated superior generalization for this architecture and dataset 4. **Adam** provided the most stable training with minimal fluctuations 5. **Architecture-optimizer interaction** matters - AlexNet responded better to SGD-style optimization

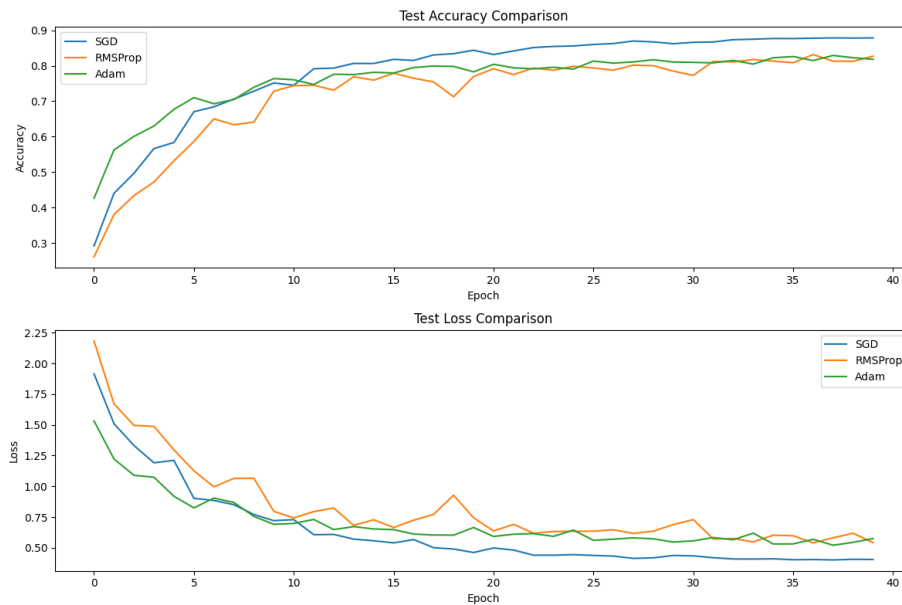


Figure 3: image

## Part (c): Performance Improvement Methods

### Method 1: Enhanced Data Augmentation


**Augmentations Applied:** - Random crop with padding - Random horizontal flip

- Random rotation ( $\pm 15^\circ$ ) - Color jitter (brightness, contrast, saturation)

**Results with ResNet18:** - **Best Test Accuracy:** ~90.2% (epoch 20) -

**Training Characteristics:** More robust training, reduced overfitting - **Im-**


**provement:** Better generalization through data diversity



## Method 2: OneCycleLR Learning Rate Policy

**Configuration:** - Max learning rate: 0.1 - 30% of training for LR increase, 70% for decrease - Cosine annealing strategy

**Results with ResNet18:** - **Best Test Accuracy:** ~95.2% (achieved faster convergence) - **Training Characteristics:** - Faster convergence to high accuracy - More efficient training with dynamic learning rate - **Best overall result** in the entire assignment



## Part (c) Analysis Summary

**Best Method:** OneCycleLR achieved the highest accuracy (95.2%)

**Key Insights:** 1. **OneCycleLR** provided the most significant improvement, achieving 95.2% accuracy 2. **Enhanced data augmentation** helped with generalization but didn't exceed baseline ResNet18 3. **Learning rate scheduling** proved more impactful than data augmentation for this dataset 4. **One Cycle policy** enables faster convergence and higher final performance

## Overall Results Summary

Model/Method	Test Accuracy	Key Characteristics
AlexNet (SGD)	85.7%	Fast training, shows overfitting
VGG16 (SGD)	90.2%	Heavy computation, stable training

Model/Method	Test Accuracy	Key Characteristics
<b>ResNet18 (SGD)</b>	<b>93.2%</b>	<b>Best architecture, stable</b>
AlexNet (RMSProp)	79.4%	Fast initial, lower final performance
AlexNet (Adam)	80.4%	Very stable, moderate performance
ResNet18 + Data Aug	90.2%	Better generalization
<b>ResNet18 + OneCycleLR</b>	<b>95.2%</b>	<b>Highest performance overall</b>




Figure 4: image