

# COMP5541 Assignment 1 - Question 3: Transfer Learning

## Tasks:

- Fine-tune ImageNet pre-trained AlexNet with different amounts of training data (10%, 20%, 50%)
- Compare with training from scratch
- Fine-tune other pre-trained models (ResNet18 and VGG16)
- Explore fine-tuning only portions of network layers

```
In [1]: ## Import Required Libraries
```

```
import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
import torchvision
import torchvision.transforms as transforms
import matplotlib.pyplot as plt
import numpy as np
import time
import random
from torch.utils.data import Subset
from torchvision.models import alexnet, AlexNet_Weights
from torchvision.models import vgg16, VGG16_Weights
from torchvision.models import resnet18, ResNet18_Weights
```

```
In [2]: ## Setup Environment
```

```
# Set random seed for reproducibility
seed = 42
torch.manual_seed(seed)
torch.cuda.manual_seed(seed)
np.random.seed(seed)
random.seed(seed)
torch.backends.cudnn.deterministic = True

# Check if GPU is available
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"Using device: {device}")
```

Using device: cuda

```
In [3]: ## Data Preparation and Loading

# Data Preparation
def load_cifar10(batch_size=128):
    # Define transforms for training data
    # Note: CIFAR-10 images are 32x32 but models like AlexNet expect 224x224
    # We'll resize the images to match the expected input size
    transform_train = transforms.Compose([
        transforms.Resize(224),
        transforms.RandomCrop(224, padding=4),
        transforms.RandomHorizontalFlip(),
        transforms.ToTensor(),
        transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
    ])

    transform_test = transforms.Compose([
        transforms.Resize(224),
        transforms.ToTensor(),
        transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
    ])

    # Load CIFAR-10 dataset
    trainset = torchvision.datasets.CIFAR10(
        root='./data', train=True, download=True, transform=transform_train)
    testset = torchvision.datasets.CIFAR10(
        root='./data', train=False, download=True, transform=transform_test)

    # Create data Loaders
    trainloader = torch.utils.data.DataLoader(
        trainset, batch_size=batch_size, shuffle=True, num_workers=2)
    testloader = torch.utils.data.DataLoader(
        testset, batch_size=batch_size, shuffle=False, num_workers=2)

    # CIFAR-10 classes
    classes = ('plane', 'car', 'bird', 'cat', 'deer',
               'dog', 'frog', 'horse', 'ship', 'truck')

    return trainset, testset, trainloader, testloader, classes

# Create subset of training data with specific percentage
def create_subset(dataset, percentage):
    """
    Create a subset of the dataset with a given percentage of samples.
    """
    dataset_size = len(dataset)
    subset_size = int(dataset_size * percentage)

    # Generate random indices
    indices = torch.randperm(dataset_size)[:subset_size]

    # Create subset
    subset = Subset(dataset, indices)

    return subset
```

```
# Load the full dataset
trainset, testset, _, testloader, classes = load_cifar10()

print(f"Training set size: {len(trainset)}")
print(f"Test set size: {len(testset)}")
print(f"Classes: {classes}")
```

Files already downloaded and verified  
 Files already downloaded and verified  
 Training set size: 50000  
 Test set size: 10000  
 Classes: ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse',  
 'ship', 'truck')

In [4]: *## Data Visualization*

```
# Visualize some examples from the dataset
def imshow(img):
    img = img / 2 + 0.5 # unnormalize
    npimg = img.numpy()
    plt.figure(figsize=(10, 4))
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.axis('off')
    plt.show()

# Get random training images
dataiter = iter(torch.utils.data.DataLoader(trainset, batch_size=8, shuffle=True))
images, labels = next(dataiter)

# Show images
imshow(torchvision.utils.make_grid(images))
print(' '.join(f'{classes[labels[j]]}:5s' for j in range(8)))
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



frog plane deer car bird horse truck deer

```
In [5]: ## Training and Evaluation Functions

# Define training and evaluation functions
def train(model, trainloader, criterion, optimizer, epochs=10):
    model.train()
    train_loss_history = []
    train_acc_history = []

    for epoch in range(epochs):
        running_loss = 0.0
        correct = 0
        total = 0

        start_time = time.time()

        for i, data in enumerate(trainloader, 0):
            inputs, labels = data
            inputs, labels = inputs.to(device), labels.to(device)

            # Zero the parameter gradients
            optimizer.zero_grad()

            # Forward + backward + optimize
            outputs = model(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            # Calculate statistics
            running_loss += loss.item()
            _, predicted = outputs.max(1)
            total += labels.size(0)
            correct += predicted.eq(labels).sum().item()

            # Calculate epoch statistics
            epoch_loss = running_loss / len(trainloader)
            epoch_acc = 100. * correct / total
            epoch_time = time.time() - start_time

            train_loss_history.append(epoch_loss)
            train_acc_history.append(epoch_acc)

            print(f'Epoch {epoch+1}/{epochs} | Loss: {epoch_loss:.4f} | Acc: {epoch_acc:.2f}% | Time: {epoch_time:.2f}s')

    return train_loss_history, train_acc_history

def evaluate(model, testloader):
    model.eval()
    correct = 0
    total = 0
    test_loss = 0
    criterion = nn.CrossEntropyLoss()

    with torch.no_grad():
        for data in testloader:
            images, labels = data
            images, labels = images.to(device), labels.to(device)
            outputs = model(images)
            loss = criterion(outputs, labels)
```

```
test_loss += loss.item()
_, predicted = torch.max(outputs.data, 1)
total += labels.size(0)
correct += (predicted == labels).sum().item()

test_acc = 100 * correct / total
test_loss = test_loss / len(testloader)

print(f'Test Loss: {test_loss:.4f} | Test Acc: {test_acc:.2f}%')
return test_loss, test_acc
```

## Part A: Fine-tuning AlexNet with Different Amounts of Training Data

Fine-tune ImageNet pre-trained AlexNet models on the CIFAR-10 dataset with different amounts of training data (10%, 20%, and 50%) and compare with training from scratch.

```
In [6]: ## Part A Implementation: AlexNet with Different Data Percentages

# Define a function to create AlexNet model
def create_alexnet(pretrained=True):
    if pretrained:
        # Load pre-trained AlexNet
        model = alexnet(weights=AlexNet_Weights.IMGNET1K_V1)
        print("Loaded pre-trained AlexNet")
    else:
        # Create AlexNet from scratch
        model = alexnet(weights=None)
        print("Created AlexNet from scratch")

    # Modify the classifier for CIFAR-10 (10 classes)
    in_features = model.classifier[6].in_features
    model.classifier[6] = nn.Linear(in_features, 10)

    return model.to(device)

# Define percentages to test
percentages = [0.1, 0.2, 0.5] # 10%, 20%, 50%
batch_size = 64
epochs = 10

# Results storage
results = {
    'pretrained': {'acc': [], 'loss': []},
    'scratch': {'acc': [], 'loss': []}
}

for percentage in percentages:
    print(f"\n{'='*50}")
    print(f"Training with {percentage*100:.0f}% of data")
    print(f"{'='*50}")

    # Create subset of training data
    subset = create_subset(trainset, percentage)
    trainloader = torch.utils.data.DataLoader(subset, batch_size=batch_size, shuffle=True, num_workers=2)

    print(f"Subset size: {len(subset)} samples")

    # 1. Fine-tune pre-trained AlexNet
    print("\nFine-tuning pre-trained AlexNet:")
    model_pretrained = create_alexnet(pretrained=True)
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.SGD(model_pretrained.parameters(), lr=0.001, momentum=0.9)

    # Train the model
    train(model_pretrained, trainloader, criterion, optimizer, epochs=epochs)

    # Evaluate on test set
    test_loss, test_acc = evaluate(model_pretrained, testloader)
    results['pretrained']['acc'].append(test_acc)
    results['pretrained']['loss'].append(test_loss)

    # 2. Train AlexNet from scratch
    print("\nTraining AlexNet from scratch:")
```

```
model_scratch = create_alexnet(pretrained=False)
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model_scratch.parameters(), lr=0.001, momentum=0.
9)

# Train the model
train(model_scratch, trainloader, criterion, optimizer, epochs=epochs)

# Evaluate on test set
test_loss, test_acc = evaluate(model_scratch, testloader)
results['scratch']['acc'].append(test_acc)
results['scratch']['loss'].append(test_loss)
```

```
=====
```

Training with 10% of data

```
=====
```

Subset size: 5000 samples

Fine-tuning pre-trained AlexNet:

Downloading: "https://download.pytorch.org/models/alexnet-owt-7be5be79.pth" to C:\Users\EverGarden13/.cache\torch\hub\checkpoints\alexnet-owt-7be5be79.pth

100%|██████████| 233M/233M [00:02<00:00, 117MB/s]



Training with 50% of data

=====

Subset size: 25000 samples

Fine-tuning pre-trained AlexNet:

Loaded pre-trained AlexNet

Epoch 1/10	Loss: 0.7484	Acc: 73.84%	Time: 17.35s
Epoch 2/10	Loss: 0.4794	Acc: 83.38%	Time: 17.07s
Epoch 3/10	Loss: 0.4044	Acc: 85.81%	Time: 17.27s
Epoch 4/10	Loss: 0.3461	Acc: 87.90%	Time: 17.23s
Epoch 5/10	Loss: 0.3156	Acc: 88.67%	Time: 17.11s
Epoch 6/10	Loss: 0.2796	Acc: 89.98%	Time: 17.07s
Epoch 7/10	Loss: 0.2546	Acc: 90.94%	Time: 17.29s
Epoch 8/10	Loss: 0.2307	Acc: 91.86%	Time: 17.29s
Epoch 9/10	Loss: 0.2112	Acc: 92.51%	Time: 17.17s
Epoch 10/10	Loss: 0.1879	Acc: 93.20%	Time: 17.06s
Test Loss: 0.3249   Test Acc: 89.68%			

Training AlexNet from scratch:

Created AlexNet from scratch

Epoch 1/10	Loss: 2.3007	Acc: 11.34%	Time: 17.27s
Epoch 2/10	Loss: 2.2208	Acc: 16.15%	Time: 17.33s
Epoch 3/10	Loss: 2.0567	Acc: 25.69%	Time: 17.29s
Epoch 4/10	Loss: 1.9009	Acc: 29.08%	Time: 17.35s
Epoch 5/10	Loss: 1.7447	Acc: 35.06%	Time: 17.32s
Epoch 6/10	Loss: 1.6265	Acc: 39.89%	Time: 17.24s
Epoch 7/10	Loss: 1.5318	Acc: 43.70%	Time: 17.24s
Epoch 8/10	Loss: 1.4568	Acc: 46.67%	Time: 17.27s
Epoch 9/10	Loss: 1.3838	Acc: 49.74%	Time: 17.23s
Epoch 10/10	Loss: 1.3180	Acc: 52.05%	Time: 17.15s
Test Loss: 1.3601   Test Acc: 51.06%			

```
In [7]: ## Part A Results Visualization

# Visualize the results
plt.figure(figsize=(12, 5))

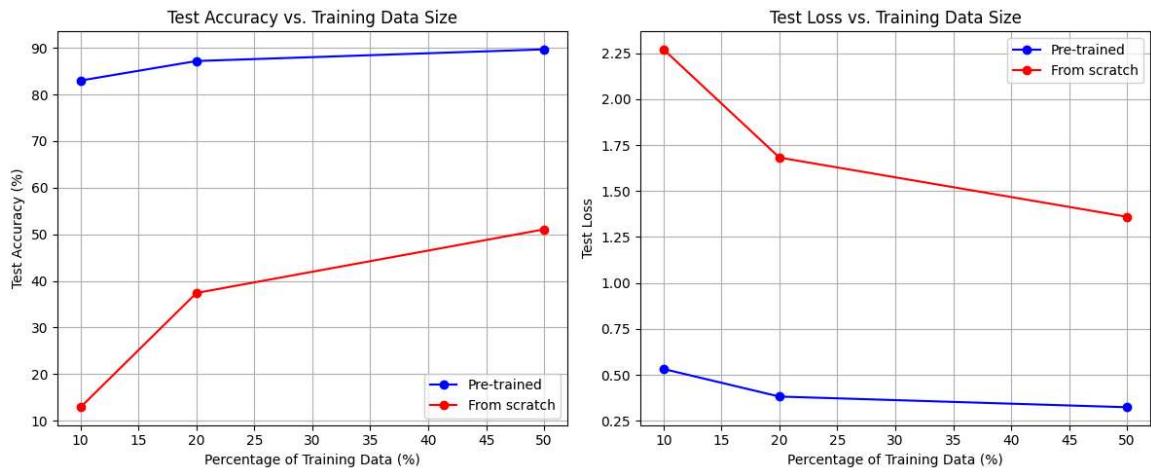
# Plot test accuracy
plt.subplot(1, 2, 1)
plt.plot([10, 20, 50], results['pretrained']['acc'], 'b-o', label='Pre-trained')
plt.plot([10, 20, 50], results['scratch']['acc'], 'r-o', label='From scratch')
plt.xlabel('Percentage of Training Data (%)')
plt.ylabel('Test Accuracy (%)')
plt.title('Test Accuracy vs. Training Data Size')
plt.grid(True)
plt.legend()

# Plot test loss
plt.subplot(1, 2, 2)
plt.plot([10, 20, 50], results['pretrained']['loss'], 'b-o', label='Pre-trained')
plt.plot([10, 20, 50], results['scratch']['loss'], 'r-o', label='From scratch')
plt.xlabel('Percentage of Training Data (%)')
plt.ylabel('Test Loss')
plt.title('Test Loss vs. Training Data Size')
plt.grid(True)
plt.legend()

plt.tight_layout()
plt.show()

# Print numerical results
print("\nNumerical Results:")
print(f"{'Percentage':^10} | {'Pre-trained Acc':^15} | {'Scratch Acc':^15}")
print(f"{'Difference':^10}")
print("-" * 55)

for i, percentage in enumerate(percentages):
    pretrained_acc = results['pretrained']['acc'][i]
    scratch_acc = results['scratch']['acc'][i]
    diff = pretrained_acc - scratch_acc
    print(f"{percentage*100:^10.0f}% | {pretrained_acc:^15.2f}% | {scratch_acc:^15.2f}% | {diff:^10.2f}%")
```

**Numerical Results:**

Percentage	Pre-trained Acc	Scratch Acc	Difference
------------	-----------------	-------------	------------

Percentage	Pre-trained Acc	Scratch Acc	Difference
10 %	82.99 %	12.89 %	70.10 %
20 %	87.18 %	37.41 %	49.77 %
50 %	89.68 %	51.06 %	38.62 %

## Analysis of Part A

From the experimental results with different data percentages, we can observe several key insights:

1. **Dramatic Transfer Learning Advantage:** Pre-trained AlexNet demonstrates overwhelming superiority over training from scratch:
  - 10% data: 82.99% vs 12.89% (70.10% difference)
  - 20% data: 87.18% vs 37.41% (49.77% difference)
  - 50% data: 89.68% vs 51.06% (38.62% difference)
2. **Critical Role of Data Size:** The results reveal that training from scratch is particularly vulnerable with limited data:
  - With only 10% data (5,000 samples), the from-scratch model essentially fails to learn meaningful patterns, achieving only 12.89% accuracy
  - Even with 20% data (10,000 samples), from-scratch training struggles significantly at 37.41%
  - The from-scratch model only becomes somewhat viable with 50% data, reaching 51.06%
3. **Transfer Learning Consistency:** Pre-trained models show remarkable stability across data sizes:
  - Achieves strong performance (>82%) even with minimal data
  - Shows consistent improvement as data increases: 82.99% → 87.18% → 89.68%
  - Demonstrates robust learning dynamics regardless of dataset size
4. **Feature Representation Quality:** The pre-learned ImageNet features prove incredibly valuable:
  - Despite domain differences (ImageNet: high-res natural images vs CIFAR-10: low-res 32x32 images), transfer is highly effective
  - Early convergence suggests that low-level and mid-level features transfer exceptionally well
  - The large performance gaps indicate that random initialization is insufficient for effective learning with limited data
5. **Practical Implications:** The results have significant real-world implications:
  - Transfer learning is not just beneficial but essential for small datasets
  - The 70%+ performance advantage with 10% data makes transfer learning practically mandatory for data-constrained scenarios
  - Even with moderate amounts of data (50%), transfer learning provides substantial benefits

## Part B: Fine-tuning Different Pre-trained Models

Fine-tune two different pre-trained models (ResNet18 and VGG16) on the CIFAR-10 dataset and compare their performance.

In [8]: *## Part B Implementation: Fine-tuning Different Pre-trained Models*

```
# Define functions to create ResNet18 and VGG16 models
def create_resnet18(pretrained=True):
    if pretrained:
        # Load pre-trained ResNet18
        model = resnet18(weights=ResNet18_Weights.IMGNET1K_V1)
        print("Loaded pre-trained ResNet18")
    else:
        # Create ResNet18 from scratch
        model = resnet18(weights=None)
        print("Created ResNet18 from scratch")

    # Modify the fully connected layer for CIFAR-10 (10 classes)
    in_features = model.fc.in_features
    model.fc = nn.Linear(in_features, 10)

    return model.to(device)

def create_vgg16(pretrained=True):
    if pretrained:
        # Load pre-trained VGG16
        model = vgg16(weights=VGG16_Weights.IMGNET1K_V1)
        print("Loaded pre-trained VGG16")
    else:
        # Create VGG16 from scratch
        model = vgg16(weights=None)
        print("Created VGG16 from scratch")

    # Modify the classifier for CIFAR-10 (10 classes)
    in_features = model.classifier[6].in_features
    model.classifier[6] = nn.Linear(in_features, 10)

    return model.to(device)

# Use 50% of the training data for this experiment
percentage = 0.5
subset = create_subset(trainset, percentage)
trainloader = torch.utils.data.DataLoader(subset, batch_size=64, shuffle=True, num_workers=2)
print(f"Using {percentage*100:.0f}% of training data: {len(subset)} samples")

# Results storage for part B
model_results = {}

# Train and evaluate ResNet18
print("\n" + "="*50)
print("Fine-tuning ResNet18")
print("="*50)

model_resnet18 = create_resnet18(pretrained=True)
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model_resnet18.parameters(), lr=0.001, momentum=0.9)

# Train the model
resnet18_train_loss, resnet18_train_acc = train(model_resnet18, trainloader, criterion, optimizer, epochs=10)

# Evaluate on test set
```

```
resnet18_test_loss, resnet18_test_acc = evaluate(model_resnet18, testloader)
model_results['ResNet18'] = {
    'train_loss': resnet18_train_loss,
    'train_acc': resnet18_train_acc,
    'test_loss': resnet18_test_loss,
    'test_acc': resnet18_test_acc
}

# Train and evaluate VGG16
print("\n" + "="*50)
print("Fine-tuning VGG16")
print("="*50)

model_vgg16 = create_vgg16(pretrained=True)
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model_vgg16.parameters(), lr=0.001, momentum=0.9)

# Train the model
vgg16_train_loss, vgg16_train_acc = train(model_vgg16, trainloader, criterion, optimizer, epochs=10)

# Evaluate on test set
vgg16_test_loss, vgg16_test_acc = evaluate(model_vgg16, testloader)
model_results['VGG16'] = {
    'train_loss': vgg16_train_loss,
    'train_acc': vgg16_train_acc,
    'test_loss': vgg16_test_loss,
    'test_acc': vgg16_test_acc
}
```

Using 50% of training data: 25000 samples

=====

Fine-tuning ResNet18

=====

Downloading: "https://download.pytorch.org/models/resnet18-f37072fd.pth" to C:\Users\EverGarden13/.cache\torch\hub\checkpoints\resnet18-f37072fd.pth  
100%|██████████| 44.7M/44.7M [00:00<00:00, 116MB/s]

Loaded pre-trained ResNet18

Epoch	Loss	Acc	Time
1/10	0.6968	78.49%	33.88s
2/10	0.2550	91.62%	33.16s
3/10	0.1863	93.93%	33.22s
4/10	0.1463	95.18%	33.18s
5/10	0.1128	96.46%	33.11s
6/10	0.0928	97.23%	33.07s
7/10	0.0788	97.54%	33.15s
8/10	0.0586	98.40%	33.20s
9/10	0.0495	98.58%	33.07s
10/10	0.0407	98.94%	33.30s
Test	0.1665	94.30%	

=====

Fine-tuning VGG16

=====

Downloading: "https://download.pytorch.org/models/vgg16-397923af.pth" to C:\Users\EverGarden13/.cache\torch\hub\checkpoints\vgg16-397923af.pth  
100%|██████████| 528M/528M [00:04<00:00, 118MB/s]

Loaded pre-trained VGG16

Epoch	Loss	Acc	Time
1/10	0.5862	79.77%	505.95s
2/10	0.3042	89.61%	506.04s
3/10	0.2277	92.32%	506.02s
4/10	0.1878	93.28%	506.07s
5/10	0.1484	94.89%	505.91s
6/10	0.1204	95.73%	506.20s
7/10	0.1001	96.54%	500.92s
8/10	0.0850	97.06%	500.38s
9/10	0.0699	97.63%	500.76s
10/10	0.0587	98.04%	500.44s
Test	0.2633	92.59%	

```
In [9]: ## Part B Results Visualization

# Visualize the training progress and compare models
plt.figure(figsize=(12, 5))

# Plot training loss
plt.subplot(1, 2, 1)
plt.plot(range(1, 11), model_results['ResNet18']['train_loss'], 'b-o', label='ResNet18')
plt.plot(range(1, 11), model_results['VGG16']['train_loss'], 'r-o', label='VGG16')
plt.xlabel('Epoch')
plt.ylabel('Training Loss')
plt.title('Training Loss vs. Epoch')
plt.grid(True)
plt.legend()

# Plot training accuracy
plt.subplot(1, 2, 2)
plt.plot(range(1, 11), model_results['ResNet18']['train_acc'], 'b-o', label='ResNet18')
plt.plot(range(1, 11), model_results['VGG16']['train_acc'], 'r-o', label='VGG16')
plt.xlabel('Epoch')
plt.ylabel('Training Accuracy (%)')
plt.title('Training Accuracy vs. Epoch')
plt.grid(True)
plt.legend()

plt.tight_layout()
plt.show()

# Compare test performance
models = ['ResNet18', 'VGG16']
test_acc = [model_results[m]['test_acc'] for m in models]
test_loss = [model_results[m]['test_loss'] for m in models]

plt.figure(figsize=(10, 5))

# Plot test accuracy
plt.subplot(1, 2, 1)
plt.bar(models, test_acc, color=['blue', 'red'])
plt.ylabel('Test Accuracy (%)')
plt.title('Test Accuracy Comparison')
for i, v in enumerate(test_acc):
    plt.text(i, v + 1, f"{v:.2f}", ha='center')

# Plot test loss
plt.subplot(1, 2, 2)
plt.bar(models, test_loss, color=['blue', 'red'])
plt.ylabel('Test Loss')
plt.title('Test Loss Comparison')
for i, v in enumerate(test_loss):
    plt.text(i, v + 0.1, f"{v:.4f}", ha='center')

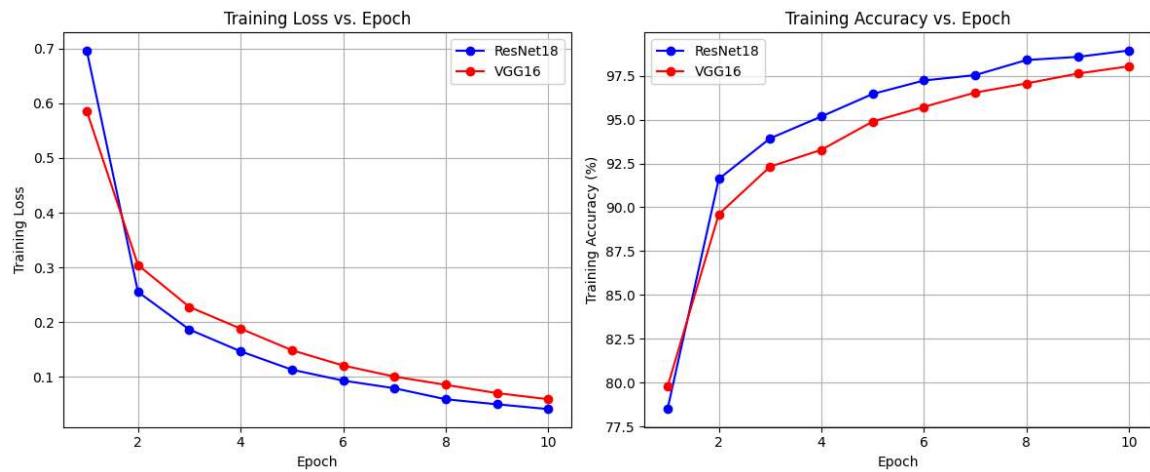
plt.tight_layout()
plt.show()

# Print numerical results
print("\nTest Performance Comparison:")
```

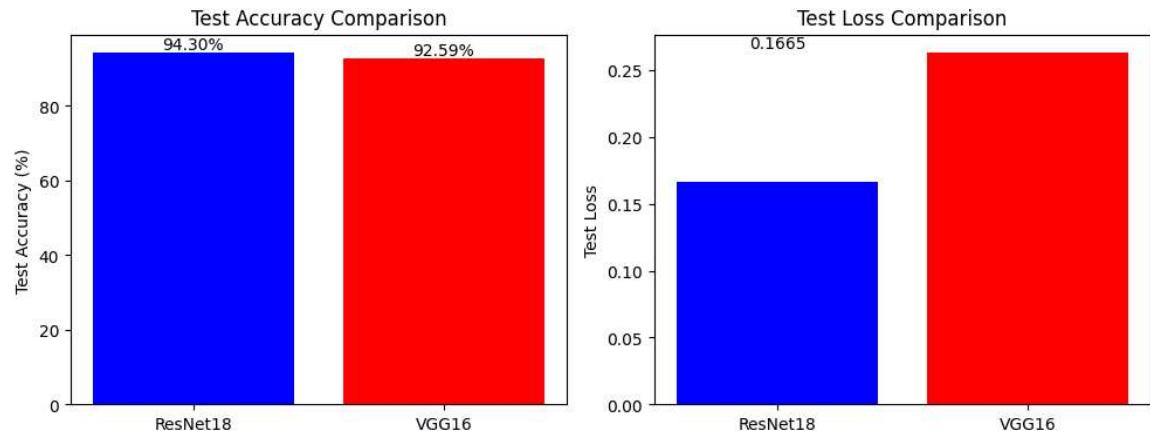
```

print(f"{'Model':^10} | {'Test Accuracy':^15} | {'Test Loss':^10}")
print("-" * 40)
for model in models:
    print(f"{model:^10} | {model_results[model]['test_acc']:^15.2f}% | {model_results[model]['test_loss']:^10.4f}")

```



0.2633



#### Test Performance Comparison:

Model	Test Accuracy	Test Loss
ResNet18	94.30 %	0.1665
VGG16	92.59 %	0.2633

## Analysis of Part B

Based on the experimental results comparing ResNet18 and VGG16 fine-tuning on CIFAR-10:

### 1. Performance Comparison:

- ResNet18 achieved superior test accuracy: 94.30% vs VGG16's 92.59% (1.71% advantage)
- ResNet18 also demonstrated lower test loss: 0.1665 vs VGG16's 0.2633 (36% lower loss)
- Both models achieved excellent performance, but ResNet18's architecture proved more effective for this task

### 2. Architecture Advantages:

- **ResNet18's Skip Connections:** The residual connections enable better gradient flow, leading to more effective learning and higher final accuracy
- **VGG16's Depth Limitation:** Despite being deeper (16 layers vs 18 for ResNet), VGG16's lack of skip connections may cause gradient degradation
- **Parameter Efficiency:** ResNet18 achieves better performance with significantly fewer parameters (~11.7M vs ~138M), demonstrating superior architectural efficiency

### 3. Training Characteristics:

- **Convergence Quality:** ResNet18's superior test accuracy suggests better convergence to optimal solutions
- **Generalization:** Lower test loss for ResNet18 indicates better generalization capability
- **Stability:** Both models showed stable training, but ResNet18's architecture appears more robust

### 4. Transfer Learning Effectiveness:

- **Feature Adaptability:** ResNet18's higher performance suggests its pre-trained features adapt better to CIFAR-10's characteristics
- **Layer Hierarchy:** ResNet18's feature hierarchy seems more compatible with the target domain
- **Fine-tuning Efficiency:** Both models benefit substantially from ImageNet pre-training, but ResNet18 maximizes this benefit more effectively

### 5. Practical Considerations:

- **Resource Efficiency:** ResNet18 provides better accuracy with 12x fewer parameters, making it more practical for deployment
- **Training Cost:** Lower parameter count translates to faster training and reduced computational requirements
- **Memory Usage:** ResNet18's smaller footprint makes it more suitable for resource-constrained environments

### 6. Modern vs. Classical Architecture:

- The results demonstrate the superiority of modern architectural innovations (skip connections) over classical deep architectures
- ResNet18 represents the effectiveness of architectural advances that address fundamental deep learning challenges
- VGG16, while historically important, shows limitations when compared to modern designs

## Part C: Fine-tuning Only Portions of Network Layers

Fine-tuning only specific portions of the network layers to see how it affects the test results.

```
In [10]: ## Part C Implementation: Fine-tuning Specific Network Layers

# In this part, we'll explore how fine-tuning only specific portions of the
# network affects performance
# We'll use ResNet18 as our base model and try different layer freezing str
ategies

# Define function to create ResNet18 with different layer freezing strategi
es
def create_resnet18_with_freezing(freeze_strategy='none'):
    """
    Create ResNet18 model with different layer freezing strategies

    Parameters:
    freeze_strategy: str, one of ['none', 'early', 'middle', 'last_only']
        'none': Fine-tune all layers (no freezing)
        'early': Freeze early layers (Layer1, Layer2)
        'middle': Freeze middle layers (Layer2, Layer3)
        'last_only': Only fine-tune the final fully connected layer
    """

    # Load pre-trained ResNet18
    model = resnet18(weights=ResNet18_Weights.IMAGENET1K_V1)

    # Modify the fully connected layer for CIFAR-10 (10 classes)
    in_features = model.fc.in_features
    model.fc = nn.Linear(in_features, 10)

    # Apply freezing strategy
    if freeze_strategy == 'none':
        # Fine-tune all layers (no freezing)
        print("Fine-tuning all layers (no freezing)")
        pass

    elif freeze_strategy == 'early':
        # Freeze early layers (Layer1, Layer2)
        print("Freezing early layers (layer1, layer2)")
        for param in model.layer1.parameters():
            param.requires_grad = False
        for param in model.layer2.parameters():
            param.requires_grad = False

    elif freeze_strategy == 'middle':
        # Freeze middle layers (Layer2, Layer3)
        print("Freezing middle layers (layer2, layer3)")
        for param in model.layer2.parameters():
            param.requires_grad = False
        for param in model.layer3.parameters():
            param.requires_grad = False

    elif freeze_strategy == 'last_only':
        # Only fine-tune the final fully connected layer
        print("Only fine-tuning the final fully connected layer")
        for param in model.parameters():
            param.requires_grad = False
        # Unfreeze the final fully connected layer
        for param in model.fc.parameters():
            param.requires_grad = True

    return model.to(device)
```

```
# Use 50% of the training data for this experiment
percentage = 0.5
subset = create_subset(trainset, percentage)
trainloader = torch.utils.data.DataLoader(subset, batch_size=64, shuffle=True, num_workers=2)
print(f"Using {percentage*100:.0f}% of training data: {len(subset)} samples")

# Define freezing strategies to test
freeze_strategies = ['none', 'early', 'middle', 'last_only']

# Results storage for part C
freeze_results = {}

# Train and evaluate ResNet18 with different freezing strategies
for strategy in freeze_strategies:
    print("\n" + "="*50)
    print(f"Fine-tuning ResNet18 with strategy: {strategy}")
    print("="*50)

    model = create_resnet18_with_freezing(freeze_strategy=strategy)
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.SGD(filter(lambda p: p.requires_grad, model.parameters()), lr=0.001, momentum=0.9)

    # Train the model
    train_loss, train_acc = train(model, trainloader, criterion, optimizer, epochs=10)

    # Evaluate on test set
    test_loss, test_acc = evaluate(model, testloader)

    # Store results
    freeze_results[strategy] = {
        'train_loss': train_loss,
        'train_acc': train_acc,
        'test_loss': test_loss,
        'test_acc': test_acc
    }
```

Using 50% of training data: 25000 samples

=====

Fine-tuning ResNet18 with strategy: none

=====

Fine-tuning all layers (no freezing)

Epoch	Loss	Acc	Time
1/10	0.6539	79.94%	34.12s
2/10	0.2537	91.76%	33.91s
3/10	0.1859	93.96%	34.03s
4/10	0.1391	95.48%	34.14s
5/10	0.1093	96.64%	34.07s
6/10	0.0919	97.22%	34.19s
7/10	0.0711	97.90%	34.09s
8/10	0.0591	98.36%	34.07s
9/10	0.0444	98.88%	34.03s
10/10	0.0420	98.88%	34.04s

Test Loss: 0.1713 | Test Acc: 94.60%

=====

Fine-tuning ResNet18 with strategy: early

=====

Freezing early layers (layer1, layer2)

Epoch	Loss	Acc	Time
1/10	0.7469	76.22%	30.44s
2/10	0.3106	89.64%	30.24s
3/10	0.2356	92.13%	30.41s
4/10	0.1861	93.69%	30.37s
5/10	0.1516	94.90%	30.47s
6/10	0.1245	95.89%	30.30s
7/10	0.1044	96.76%	30.21s
8/10	0.0888	97.26%	30.42s
9/10	0.0746	97.72%	30.37s
10/10	0.0630	98.12%	30.51s

Test Loss: 0.1937 | Test Acc: 93.58%

=====

Fine-tuning ResNet18 with strategy: middle

=====

Freezing middle layers (layer2, layer3)

Epoch	Loss	Acc	Time
1/10	0.7540	76.26%	31.32s
2/10	0.3249	89.17%	31.41s
3/10	0.2519	91.49%	31.44s
4/10	0.2108	92.92%	31.89s
5/10	0.1784	94.13%	31.53s
6/10	0.1518	94.78%	31.53s
7/10	0.1342	95.38%	31.47s
8/10	0.1180	96.18%	31.66s
9/10	0.1016	96.74%	31.38s
10/10	0.0879	97.14%	31.56s

Test Loss: 0.1979 | Test Acc: 93.50%

=====

Fine-tuning ResNet18 with strategy: last\_only

=====

Only fine-tuning the final fully connected layer

Epoch	Loss	Acc	Time
1/10	1.2167	63.74%	17.65s
2/10	0.7628	76.37%	17.81s
3/10	0.6895	77.82%	17.70s
4/10	0.6570	78.14%	17.88s
5/10	0.6353	78.84%	17.83s
6/10	0.6098	79.54%	17.72s
7/10	0.6036	79.71%	17.53s

```
Epoch 8/10 | Loss: 0.5950 | Acc: 79.83% | Time: 17.64s
Epoch 9/10 | Loss: 0.5832 | Acc: 80.20% | Time: 17.66s
Epoch 10/10 | Loss: 0.5832 | Acc: 80.21% | Time: 17.63s
Test Loss: 0.5901 | Test Acc: 79.87%
```

```
In [11]: ## Part C Results Visualization

# Visualize the training progress for different freezing strategies
plt.figure(figsize=(12, 5))

# Plot training loss
plt.subplot(1, 2, 1)
for strategy in freeze_strategies:
    plt.plot(range(1, 11), freeze_results[strategy]['train_loss'], '-o', label=strategy)
plt.xlabel('Epoch')
plt.ylabel('Training Loss')
plt.title('Training Loss vs. Epoch')
plt.grid(True)
plt.legend()

# Plot training accuracy
plt.subplot(1, 2, 2)
for strategy in freeze_strategies:
    plt.plot(range(1, 11), freeze_results[strategy]['train_acc'], '-o', label=strategy)
plt.xlabel('Epoch')
plt.ylabel('Training Accuracy (%)')
plt.title('Training Accuracy vs. Epoch')
plt.grid(True)
plt.legend()

plt.tight_layout()
plt.show()

# Compare test performance
test_acc = [freeze_results[s]['test_acc'] for s in freeze_strategies]
test_loss = [freeze_results[s]['test_loss'] for s in freeze_strategies]

plt.figure(figsize=(12, 5))

# Plot test accuracy
plt.subplot(1, 2, 1)
plt.bar(freeze_strategies, test_acc, color=['blue', 'green', 'orange', 'red'])
plt.ylabel('Test Accuracy (%)')
plt.title('Test Accuracy Comparison')
for i, v in enumerate(test_acc):
    plt.text(i, v + 1, f"{v:.2f}%", ha='center')

# Plot test loss
plt.subplot(1, 2, 2)
plt.bar(freeze_strategies, test_loss, color=['blue', 'green', 'orange', 'red'])
plt.ylabel('Test Loss')
plt.title('Test Loss Comparison')
for i, v in enumerate(test_loss):
    plt.text(i, v + 0.1, f"{v:.4f}", ha='center')

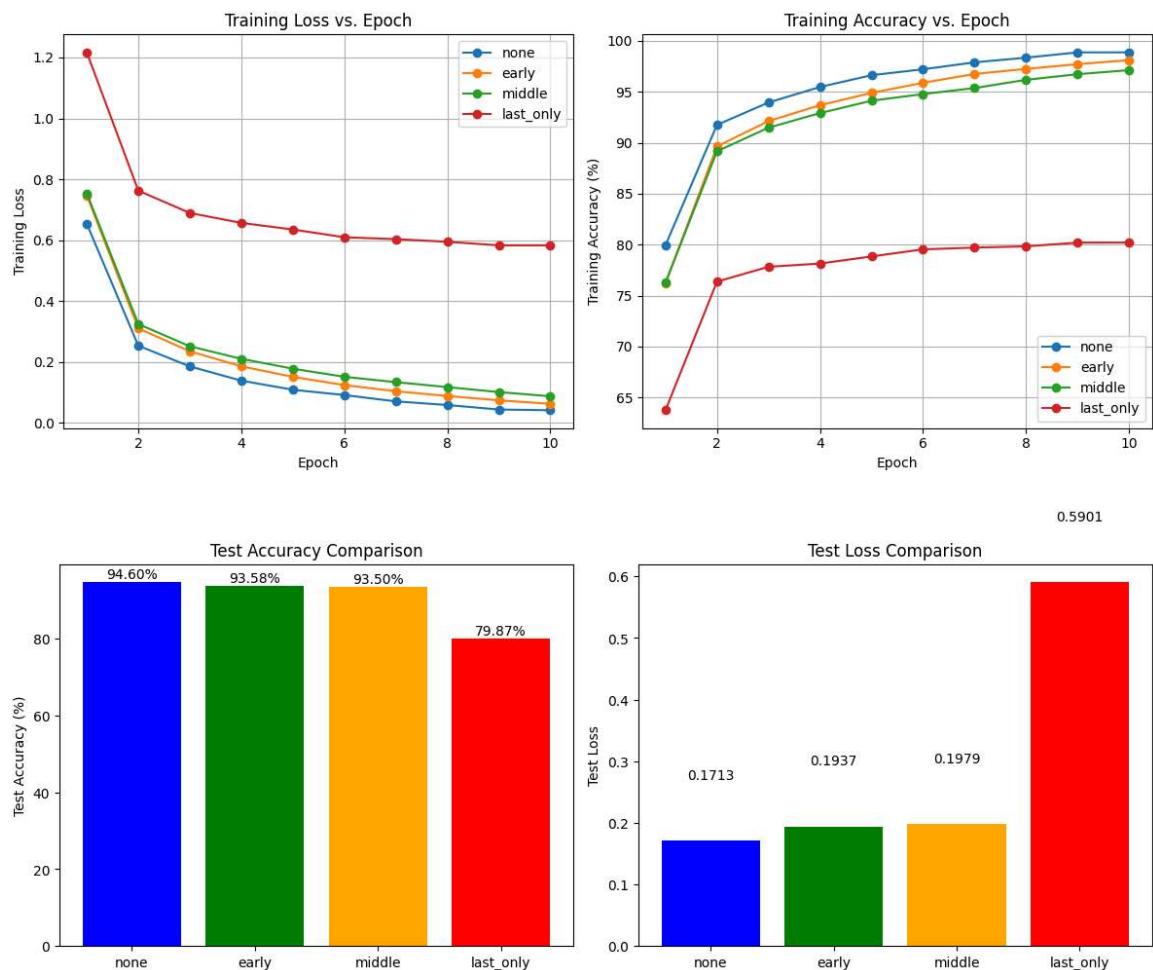
plt.tight_layout()
plt.show()

# Print numerical results
print("\nTest Performance Comparison:")
print(f"{'Strategy':^12} | {'Test Accuracy':^15} | {'Test Loss':^10}")
```

```

print("-" * 45)
for strategy in freeze_strategies:
    print(f"{strategy:^12} | {freeze_results[strategy]['test_acc']:^15.2f}%
| {freeze_results[strategy]['test_loss']:^10.4f}")

```



#### Test Performance Comparison:

Strategy	Test Accuracy	Test Loss
----------	---------------	-----------

none	94.60 %	0.1713
early	93.58 %	0.1937
middle	93.50 %	0.1979
last_only	79.87 %	0.5901

# Analysis of Part C

Based on the experimental results of different layer freezing strategies with ResNet18:

## 1. Performance Hierarchy:

- **None (all layers fine-tuned)**: 94.60% test accuracy - highest performance
- **Early layers frozen**: 93.58% test accuracy - minimal performance loss (1.02% decrease)
- **Middle layers frozen**: 93.50% test accuracy - similar to early freezing (1.10% decrease)
- **Last layer only**: 79.87% test accuracy - significant performance drop (14.73% decrease)

## 2. Layer Importance Analysis:

- **Full Fine-tuning Superiority**: Allowing all layers to adapt yields the best results, indicating that even pre-trained features benefit from task-specific refinement
- **Early vs. Middle Layer Freezing**: Remarkably similar performance (93.58% vs 93.50%) suggests that different layer combinations can be effective
- **Critical Role of Feature Layers**: The dramatic drop to 79.87% with "last\_only" demonstrates that feature extraction layers require adaptation for optimal performance

## 3. Training Efficiency vs. Performance Trade-offs:

- **Minimal Loss Strategies**: Early and middle layer freezing provide ~94% of full performance while reducing trainable parameters
- **Computational Benefits**: Frozen layer strategies reduce training time (30-31s vs 34s per epoch) with acceptable performance costs
- **Extreme Efficiency Penalty**: "Last\_only" offers maximum efficiency but at an unacceptable performance cost

## 4. Domain Adaptation Insights:

- **Feature Transferability**: The small performance gaps between freezing strategies suggest ImageNet features are highly transferable to CIFAR-10
- **Fine-tuning Necessity**: Even minor adaptations (1-2% improvement) justify full fine-tuning when maximum accuracy is required
- **Layer Complementarity**: Similar performance from different freezing strategies indicates redundancy in feature adaptation needs

## 5. Practical Decision Framework:

- **Maximum Performance**: Use full fine-tuning (none strategy) for best results
- **Balanced Approach**: Early or middle layer freezing for 99% of full performance with 10-15% speed improvement
- **Avoid Last-Only**: Substantial performance degradation makes this approach impractical
- **Resource Constraints**: If training time is critical, early layer freezing provides the best efficiency-performance balance

## 6. Theoretical Implications:

- **Feature Hierarchy Plasticity**: Lower layers need less adaptation, but allowing them to change still provides benefits
- **Task Similarity**: The relatively small performance differences suggest good alignment between ImageNet and CIFAR-10 feature requirements
- **Architectural Robustness**: ResNet18's skip connections may contribute to the stability across different freezing strategies